

Lab 1 - College Football Salary Recommendation

July 24, 2022

1 LAB 1 - Marley Akonnor

Research Question: How can we recommend the best salary (totalpay or schoolpay – your choice) for our next head football coach?

##Import Required Packages

```
[150]: #Import appropriate libraries for analysis
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
import statsmodels.api as sm
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.formula.api as smf
%matplotlib inline
```

##Import and cleaning of the datasets:

```
[89]: #Import the coaches dataset
coaches_df = pd.read_csv('/Users/m/Documents/M.S Syracuse Data Science/Courses/
↳ IST 718 - Big Data Analytics/Labs/Coaches9.csv')
coaches_df.head()
```

```
[89]:
```

	School	Conference	Coach	SchoolPay	\
0	Air Force	Mt. West	Troy Calhoun	885000	
1	Akron	MAC	Terry Bowden	\$411,000	
2	Alabama	SEC	Nick Saban	\$8,307,000	
3	Alabama at Birmingham	C-USA	Bill Clark	\$900,000	
4	Appalachian State	Sun Belt	Scott Satterfield	\$712,500	

	TotalPay	Bonus	BonusPaid	AssistantPay	Buyout
0	885000	247000	--	\$0	--
1	\$412,500	\$225,000	\$50,000	\$0	\$688,500
2	\$8,307,000	\$1,100,000	\$500,000	\$0	\$33,600,000
3	\$900,000	\$950,000	\$165,471	\$0	\$3,847,500

4	\$712,500	\$295,000	\$145,000	\$0	\$2,160,417
---	-----------	-----------	-----------	-----	-------------

```
[90]: #Clean the dataset by removing superfluous columns
coaches_df = coaches_df.drop(columns=['Bonus', 'Buyout', 'AssistantPay'])
```

```
[91]: #Continues data cleaning:
#Utilize a function by Allen Flores to expedite data cleaning.
#This function removes commas, dollar signs, and extra spaces. Use this_
→function on the appropriate columns
```

```
def convert2num(data_column):
    col = data_column.str.replace(',', '')
    col = col.str.replace('$', '')
    col = col.str.replace(" ", "")
    col = col.str.replace("--", "0")
    col = col.fillna(0)
    col = col.astype(float)
    return col

coaches_df['SchoolPay'] = convert2num(coaches_df['SchoolPay'])
coaches_df['TotalPay'] = convert2num(coaches_df['TotalPay'])
coaches_df['BonusPaid'] = convert2num(coaches_df['BonusPaid'])

coaches_df.head()
```

<ipython-input-91-33db4cb2753b>:7: FutureWarning: The default value of regex will change from True to False in a future version. In addition, single character regular expressions will*not* be treated as literal strings when regex=True.

```
col = col.str.replace('$', '')
```

```
[91]:
```

	School	Conference	Coach	SchoolPay	TotalPay \
0	Air Force	Mt. West	Troy Calhoun	885000.0	885000.0
1	Akron	MAC	Terry Bowden	411000.0	412500.0
2	Alabama	SEC	Nick Saban	8307000.0	8307000.0
3	Alabama at Birmingham	C-USA	Bill Clark	900000.0	900000.0
4	Appalachian State	Sun Belt	Scott Satterfield	712500.0	712500.0

	BonusPaid
0	0.0
1	50000.0
2	500000.0
3	165471.0
4	145000.0

```
[92]: #Import the stadium dataset
#Use the converter function on the Capacity column
#Change the "College" column name to "School" in preparation for the merge
```

```

stadium_df = pd.read_csv('/Users/m/Documents/M.S Syracuse Data Science/Courses/
↳IST 718 - Big Data Analytics/Labs/Stadium_Size.csv')
stadium_df['Capacity'] = convert2num(stadium_df['Capacity'])
stadium_df.rename(columns={'College':'School'}, inplace=True)
stadium_df.head(5)

```

<ipython-input-91-33db4cb2753b>:7: FutureWarning: The default value of regex will change from True to False in a future version. In addition, single character regular expressions will*not* be treated as literal strings when regex=True.

```
col = col.str.replace('$', '')
```

```
[92]:
```

	Stadium	School	Conference	Capacity	Opened
0	Michigan Stadium	Michigan	Big Ten	107601.0	1927
1	Beaver Stadium	Penn State	Big Ten	106572.0	1960
2	Ohio Stadium	Ohio State	Big Ten	104944.0	1922
3	Kyle Field	Texas A&M	SEC	102733.0	1904
4	Neyland Stadium	Tennessee	SEC	102521.0	1921

```
[93]: #Import the revenue dataset
#Use the converter function on the "Revenue" and "Expenses" column
#Rename "RK" to "Rank" for clarity
revenue_df = pd.read_csv('/Users/m/Documents/M.S Syracuse Data Science/Courses/
↳IST 718 - Big Data Analytics/Labs/revenue.csv')
revenue_df['Revenue'] = convert2num(revenue_df['Revenue'])
revenue_df['Expenses'] = convert2num(revenue_df['Expenses'])
revenue_df.rename(columns={'RK':'Rank'}, inplace=True)
revenue_df.head(5)

```

<ipython-input-91-33db4cb2753b>:7: FutureWarning: The default value of regex will change from True to False in a future version. In addition, single character regular expressions will*not* be treated as literal strings when regex=True.

```
col = col.str.replace('$', '')
```

```
[93]:
```

	Rank	School	Conference	Revenue	Expenses
0	1	Texas	Big 12	223879781.0	204234897.0
1	2	Texas A&M	SEC	212748002.0	169012456.0
2	3	Ohio State	Big Ten	210548239.0	220572956.0
3	4	Michigan	Big Ten	197820410.0	190952175.0
4	5	Georgia	SEC	174042482.0	143299554.0

```
[94]: #Import the Graduation Rates dataset
#Use the converter function on the "gsr" and "fgr" column
#Change the "school" column name to "School" in preparation for the merge
gradrate_df = pd.read_csv('/Users/m/Documents/M.S Syracuse Data Science/Courses/
↳IST 718 - Big Data Analytics/Labs/grad_rates.csv')
gradrate_df.rename(columns={'school':'School'}, inplace=True)

```

```
gradrate_df.head(10)
```

```
[94]:
```

	x	School	conf	sport	state	\
0	2012	Abilene Christian	Southland Conference	Football	TX	
1	2012	Akron	Mid-American Conference	Football	OH	
2	2012	Alabama A&M	Southwestern Athletic Conf.	Football	AL	
3	2012	Alabama State	Southwestern Athletic Conf.	Football	AL	
4	2012	Alabama	Southeastern Conference	Football	AL	
5	2012	Alabama at Birmingham	Conference USA	Football	AL	
6	2012	University at Albany	Colonial Athletic Association	Football	NY	
7	2012	Alcorn State	Southwestern Athletic Conf.	Football	MS	
8	2012	Appalachian State	Sun Belt Conference	Football	NC	
9	2012	Arizona State	Pac-12 Conference	Football	AZ	

	gsr	fgr	Unnamed: 7
0	70	47.0	NaN
1	75	61.0	NaN
2	59	49.0	NaN
3	58	39.0	NaN
4	85	65.0	NaN
5	71	51.0	NaN
6	88	63.0	NaN
7	58	40.0	NaN
8	75	67.0	NaN
9	75	60.0	NaN

```
[95]: #Merge the dataframes together
#Begin by merging the Coaches data with the Stadium data and then merge Revenue
→and the grauation rate
mergeUno = coaches_df.merge(stadium_df, on='School', how='left')
mergeDos = mergeUno.merge(revenue_df, on='School', how='left')
coaches_comp = mergeDos.merge(gradrate_df, on='School', how='left')
coaches_comp.head(10)
```

```
[95]:
```

	School	Conference_x	Coach	SchoolPay	\
0	Air Force	Mt. West	Troy Calhoun	885000.0	
1	Akron	MAC	Terry Bowden	411000.0	
2	Alabama	SEC	Nick Saban	8307000.0	
3	Alabama at Birmingham	C-USA	Bill Clark	900000.0	
4	Appalachian State	Sun Belt	Scott Satterfield	712500.0	
5	Arizona	Pac-12	Kevin Sumlin	1600000.0	
6	Arizona State	Pac-12	Herm Edwards	2000000.0	
7	Arizona State	Pac-12	Herm Edwards	2000000.0	
8	Arkansas	SEC	Chad Morris	3500000.0	
9	Arkansas State	Sun Belt	Blake Anderson	825000.0	

	TotalPay	BonusPaid	Stadium	Conference_y	Capacity	\
--	----------	-----------	---------	--------------	----------	---

0	885000.0	0.0	Falcon Stadium	Mountain West	52237.0
1	412500.0	50000.0	InfoCision Stadium	MAC	30000.0
2	8307000.0	500000.0	Bryant Denny Stadium	SEC	101821.0
3	900000.0	165471.0	NaN	NaN	NaN
4	712500.0	145000.0	Kidd Brewer Stadium	Sun Belt	24150.0
5	2000000.0	0.0	Arizona Stadium	Pac 12	56037.0
6	2000000.0	0.0	Sun Devil Stadium	Pac 12	56232.0
7	2000000.0	0.0	Sun Devil Stadium	Pac 12	56232.0
8	3500000.0	0.0	Razorback Stadium	SEC	72000.0
9	825000.0	25000.0	Centennial Bank Stadium	Sun Belt	30964.0

	Opened	...	Conference	Revenue	Expenses	x	\
0	1962.0	...	Mt. West	60009782.0	54192115.0	2012.0	
1	2009.0	...	MAC	37194485.0	37275978.0	2012.0	
2	1929.0	...	SEC	164090889.0	185317681.0	2012.0	
3	NaN	...	C-USA	36494381.0	35697117.0	2012.0	
4	1962.0	...	Sun Belt	37996512.0	37773447.0	2012.0	
5	1928.0	...	Pac-12	105091389.0	100565835.0	2012.0	
6	1958.0	...	Pac-12	121698840.0	118404377.0	2012.0	
7	1958.0	...	Pac-12	121698840.0	118404377.0	2012.0	
8	1938.0	...	SEC	137497788.0	129620361.0	2012.0	
9	2002.0	...	Sun Belt	35550592.0	35550592.0	NaN	

		conf	sport	state	gsr	fgr	Unnamed: 7
0	Mountain West Conference	Football	CO	80.0	NaN		NaN
1	Mid-American Conference	Football	OH	75.0	61.0		NaN
2	Southeastern Conference	Football	AL	85.0	65.0		NaN
3	Conference USA	Football	AL	71.0	51.0		NaN
4	Sun Belt Conference	Football	NC	75.0	67.0		NaN
5	Pac-12 Conference	Football	AZ	76.0	58.0		NaN
6	Pac-12 Conference	Football	AZ	75.0	60.0		NaN
7	Sun Belt Conference	Football	AR	80.0	62.0		NaN
8	Southeastern Conference	Football	AR	67.0	47.0		NaN
9	NaN	NaN	NaN	NaN	NaN		NaN

[10 rows x 21 columns]

```
[96]: #Next remove extraneous columns
coaches_comp = coaches_comp.drop(['Stadium','Conference_y', 'Opened',
    → 'Conference', 'x', 'conf', 'sport', 'Unnamed: 7'], axis=1)

#Rename conference
coaches_comp.rename(columns={'Conference_x': 'Conference'}, inplace=True)
coaches_comp.head(10)
```

```
[96]:          School Conference          Coach SchoolPay TotalPay \
0          Air Force  Mt. West      Troy Calhoun  885000.0  885000.0
```

1	Akron	MAC	Terry Bowden	411000.0	412500.0
2	Alabama	SEC	Nick Saban	8307000.0	8307000.0
3	Alabama at Birmingham	C-USA	Bill Clark	900000.0	900000.0
4	Appalachian State	Sun Belt	Scott Satterfield	712500.0	712500.0
5	Arizona	Pac-12	Kevin Sumlin	1600000.0	2000000.0
6	Arizona State	Pac-12	Herm Edwards	2000000.0	2000000.0
7	Arizona State	Pac-12	Herm Edwards	2000000.0	2000000.0
8	Arkansas	SEC	Chad Morris	3500000.0	3500000.0
9	Arkansas State	Sun Belt	Blake Anderson	825000.0	825000.0

	BonusPaid	Capacity	Rank	Revenue	Expenses	state	gsr	fgr
0	0.0	52237.0	57.0	60009782.0	54192115.0	CO	80.0	NaN
1	50000.0	30000.0	84.0	37194485.0	37275978.0	OH	75.0	61.0
2	500000.0	101821.0	7.0	164090889.0	185317681.0	AL	85.0	65.0
3	165471.0	NaN	86.0	36494381.0	35697117.0	AL	71.0	51.0
4	145000.0	24150.0	81.0	37996512.0	37773447.0	NC	75.0	67.0
5	0.0	56037.0	38.0	105091389.0	100565835.0	AZ	76.0	58.0
6	0.0	56232.0	27.0	121698840.0	118404377.0	AZ	75.0	60.0
7	0.0	56232.0	27.0	121698840.0	118404377.0	AR	80.0	62.0
8	0.0	72000.0	20.0	137497788.0	129620361.0	AR	67.0	47.0
9	25000.0	30964.0	92.0	35550592.0	35550592.0	NaN	NaN	NaN

```
[97]: #Get an understanding of what types are in the dataset for classification,
      ↪linear or logistic regression
coaches_comp.dtypes
```

```
[97]: School      object
Conference  object
Coach       object
SchoolPay   float64
TotalPay    float64
BonusPaid   float64
Capacity    float64
Rank        float64
Revenue     float64
Expenses    float64
state       object
gsr         float64
fgr         float64
dtype: object
```

```
[98]: #Convert "Conference" and "Rank" to categorical
coaches_comp['Conference'] = coaches_comp.Conference.astype('category')
coaches_comp['Rank'] = coaches_comp.Rank.astype('category')
```

```
##Check for missing values:
```

```
[99]: #Check for null values in the dataframe
null_check = coaches_comp.isnull().values.any()
tot_null = coaches_comp.isnull().sum()
print(null_check)
print(tot_null)
```

```
True
School      0
Conference   0
Coach        0
SchoolPay    0
TotalPay     0
BonusPaid    0
Capacity    15
Rank         20
Revenue      20
Expenses     20
state        5
gsr          5
fgr          7
dtype: int64
```

```
[100]: #Colleges missing TotalPay
coaches_comp.loc[coaches_comp['TotalPay'] == 0]
```

```
[100]:
```

	School	Conference	Coach	SchoolPay	TotalPay \
13	Baylor	Big 12	Matt Rhule	0.0	0.0
17	Brigham Young	Ind.	Kalani Sitake	0.0	0.0
92	Rice	C-USA	Mike Bloomgren	0.0	0.0
100	Southern Methodist	AAC	Sonny Dykes	0.0	0.0

	BonusPaid	Capacity	Rank	Revenue	Expenses	state	gsr	fgr
13	0.0	45140.0	NaN	NaN	NaN	TX	90.0	61.0
17	0.0	NaN	NaN	NaN	NaN	UT	64.0	43.0
92	0.0	47000.0	NaN	NaN	NaN	TX	87.0	80.0
100	0.0	NaN	NaN	NaN	NaN	TX	72.0	63.0

```
[101]: #Colleges missing Stadium Capacity
coaches_comp[coaches_comp['Capacity'].isna()]
```

```
[101]:
```

	School	Conference	Coach	SchoolPay	TotalPay \
3	Alabama at Birmingham	C-USA	Bill Clark	900000.0	900000.0
17	Brigham Young	Ind.	Kalani Sitake	0.0	0.0
22	Charlotte	C-USA	Brad Lambert	625000.0	625000.0
51	Liberty	Ind.	Turner Gill	947281.0	947281.0
59	Massachusetts	Ind.	Mark Whipple	500000.0	500000.0
61	Miami (Fla.)	ACC	Mark Richt	4058061.0	4058061.0
62	Miami (Ohio)	MAC	Chuck Martin	524826.0	524826.0

73	Nevada-Las Vegas	Mt. West	Tony Sanchez	600000.0	600000.0
77	North Carolina State	ACC	Dave Doeren	3000000.0	3000000.0
99	Southern California	Pac-12	Clay Helton	2625965.0	2625965.0
100	Southern Methodist	AAC	Sonny Dykes	0.0	0.0
101	Southern Mississippi	C-USA	Jay Hopson	500000.0	500000.0
107	Texas Christian	Big 12	Gary Patterson	4840717.0	4840717.0
110	Texas-El Paso	C-USA	Dana Dimel	700000.0	700000.0
111	Texas-San Antonio	C-USA	Frank Wilson	1100000.0	1100000.0

	BonusPaid	Capacity	Rank	Revenue	Expenses	state	gsr	fgr
3	165471.0	NaN	86.0	36494381.0	35697117.0	AL	71.0	51.0
17	0.0	NaN	NaN	NaN	NaN	UT	64.0	43.0
22	0.0	NaN	76.0	40866430.0	38454232.0	NC	80.0	NaN
51	0.0	NaN	NaN	NaN	NaN	VA	67.0	65.0
59	20000.0	NaN	67.0	49461013.0	48445234.0	MA	68.0	51.0
61	0.0	NaN	NaN	NaN	NaN	FL	81.0	50.0
62	12669.0	NaN	79.0	39644248.0	38650630.0	OH	89.0	77.0
73	0.0	NaN	64.0	50784275.0	50445250.0	NV	63.0	60.0
77	275000.0	NaN	47.0	92724548.0	90100025.0	NC	71.0	54.0
99	0.0	NaN	NaN	NaN	NaN	CA	82.0	62.0
100	0.0	NaN	NaN	NaN	NaN	TX	72.0	63.0
101	15000.0	NaN	125.0	25687189.0	25687189.0	MS	72.0	58.0
107	0.0	NaN	NaN	NaN	NaN	TX	67.0	52.0
110	0.0	NaN	99.0	33481654.0	32980563.0	TX	85.0	67.0
111	17500.0	NaN	97.0	33657739.0	32278508.0	TX	81.0	59.0

```
[102]: #Colleges missing Revenue
coaches_comp[coaches_comp['Revenue'].isna()]
```

```
[102]:
```

	School	Conference	Coach	SchoolPay	TotalPay \
10	Army	Ind.	Jeff Monken	932521.0	932521.0
13	Baylor	Big 12	Matt Rhule	0.0	0.0
15	Boston College	ACC	Steve Addazio	2514859.0	2514859.0
17	Brigham Young	Ind.	Kalani Sitake	0.0	0.0
29	Duke	ACC	David Cutcliffe	2540928.0	2540928.0
51	Liberty	Ind.	Turner Gill	947281.0	947281.0
61	Miami (Fla.)	ACC	Mark Richt	4058061.0	4058061.0
70	Navy	AAC	Ken Niumatalolo	2163000.0	2163000.0
80	Northwestern	Big Ten	Pat Fitzgerald	3619775.0	3619775.0
81	Notre Dame	Ind.	Brian Kelly	2129638.0	2129638.0
90	Pittsburgh	ACC	Pat Narduzzi	3017718.0	3017718.0
92	Rice	C-USA	Mike Bloomgren	0.0	0.0
99	Southern California	Pac-12	Clay Helton	2625965.0	2625965.0
100	Southern Methodist	AAC	Sonny Dykes	0.0	0.0
102	Stanford	Pac-12	David Shaw	4311543.0	4311543.0
107	Texas Christian	Big 12	Gary Patterson	4840717.0	4840717.0
114	Tulane	AAC	Willie Fritz	1629000.0	1629000.0

115	Tulsa	AAC	Philip Montgomery	1518177.0	1518177.0
119	Vanderbilt	SEC	Derek Mason	2812523.0	2812523.0
122	Wake Forest	ACC	Dave Clawson	1831580.0	1831580.0

	BonusPaid	Capacity	Rank	Revenue	Expenses	state	gsr	fgr
10	0.0	38000.0	NaN	NaN	NaN	NaN	NaN	NaN
13	0.0	45140.0	NaN	NaN	NaN	TX	90.0	61.0
15	0.0	44500.0	NaN	NaN	NaN	MA	90.0	81.0
17	0.0	NaN	NaN	NaN	NaN	UT	64.0	43.0
29	0.0	40004.0	NaN	NaN	NaN	NC	96.0	90.0
51	0.0	NaN	NaN	NaN	NaN	VA	67.0	65.0
61	0.0	NaN	NaN	NaN	NaN	FL	81.0	50.0
70	0.0	34000.0	NaN	NaN	NaN	NaN	NaN	NaN
80	0.0	47130.0	NaN	NaN	NaN	IL	100.0	93.0
81	0.0	77622.0	NaN	NaN	NaN	IN	92.0	84.0
90	0.0	65500.0	NaN	NaN	NaN	PA	78.0	47.0
92	0.0	47000.0	NaN	NaN	NaN	TX	87.0	80.0
99	0.0	NaN	NaN	NaN	NaN	CA	82.0	62.0
100	0.0	NaN	NaN	NaN	NaN	TX	72.0	63.0
102	0.0	50000.0	NaN	NaN	NaN	CA	91.0	88.0
107	0.0	NaN	NaN	NaN	NaN	TX	67.0	52.0
114	0.0	30000.0	NaN	NaN	NaN	LA	89.0	73.0
115	0.0	35542.0	NaN	NaN	NaN	OK	78.0	72.0
119	0.0	40550.0	NaN	NaN	NaN	TN	95.0	81.0
122	0.0	31500.0	NaN	NaN	NaN	NC	86.0	73.0

```
[103]: #Colleges missing Expenses
coaches_comp[coaches_comp['Expenses'].isna()]
```

```
[103]:
```

	School	Conference	Coach	SchoolPay	TotalPay \
10	Army	Ind.	Jeff Monken	932521.0	932521.0
13	Baylor	Big 12	Matt Rhule	0.0	0.0
15	Boston College	ACC	Steve Addazio	2514859.0	2514859.0
17	Brigham Young	Ind.	Kalani Sitake	0.0	0.0
29	Duke	ACC	David Cutcliffe	2540928.0	2540928.0
51	Liberty	Ind.	Turner Gill	947281.0	947281.0
61	Miami (Fla.)	ACC	Mark Richt	4058061.0	4058061.0
70	Navy	AAC	Ken Niumatalolo	2163000.0	2163000.0
80	Northwestern	Big Ten	Pat Fitzgerald	3619775.0	3619775.0
81	Notre Dame	Ind.	Brian Kelly	2129638.0	2129638.0
90	Pittsburgh	ACC	Pat Narduzzi	3017718.0	3017718.0
92	Rice	C-USA	Mike Bloomgren	0.0	0.0
99	Southern California	Pac-12	Clay Helton	2625965.0	2625965.0
100	Southern Methodist	AAC	Sonny Dykes	0.0	0.0
102	Stanford	Pac-12	David Shaw	4311543.0	4311543.0
107	Texas Christian	Big 12	Gary Patterson	4840717.0	4840717.0
114	Tulane	AAC	Willie Fritz	1629000.0	1629000.0

115	Tulsa	AAC	Philip Montgomery	1518177.0	1518177.0
119	Vanderbilt	SEC	Derek Mason	2812523.0	2812523.0
122	Wake Forest	ACC	Dave Clawson	1831580.0	1831580.0

	BonusPaid	Capacity	Rank	Revenue	Expenses	state	gsr	fgr
10	0.0	38000.0	NaN	NaN	NaN	NaN	NaN	NaN
13	0.0	45140.0	NaN	NaN	NaN	TX	90.0	61.0
15	0.0	44500.0	NaN	NaN	NaN	MA	90.0	81.0
17	0.0	NaN	NaN	NaN	NaN	UT	64.0	43.0
29	0.0	40004.0	NaN	NaN	NaN	NC	96.0	90.0
51	0.0	NaN	NaN	NaN	NaN	VA	67.0	65.0
61	0.0	NaN	NaN	NaN	NaN	FL	81.0	50.0
70	0.0	34000.0	NaN	NaN	NaN	NaN	NaN	NaN
80	0.0	47130.0	NaN	NaN	NaN	IL	100.0	93.0
81	0.0	77622.0	NaN	NaN	NaN	IN	92.0	84.0
90	0.0	65500.0	NaN	NaN	NaN	PA	78.0	47.0
92	0.0	47000.0	NaN	NaN	NaN	TX	87.0	80.0
99	0.0	NaN	NaN	NaN	NaN	CA	82.0	62.0
100	0.0	NaN	NaN	NaN	NaN	TX	72.0	63.0
102	0.0	50000.0	NaN	NaN	NaN	CA	91.0	88.0
107	0.0	NaN	NaN	NaN	NaN	TX	67.0	52.0
114	0.0	30000.0	NaN	NaN	NaN	LA	89.0	73.0
115	0.0	35542.0	NaN	NaN	NaN	OK	78.0	72.0
119	0.0	40550.0	NaN	NaN	NaN	TN	95.0	81.0
122	0.0	31500.0	NaN	NaN	NaN	NC	86.0	73.0

```
[104]: #Colleges missing Federal Graduation Rate
coaches_comp[coaches_comp['fgr'].isna()]
```

```
[104]:
```

	School	Conference	Coach	SchoolPay	TotalPay \
0	Air Force	Mt. West	Troy Calhoun	885000.0	885000.0
9	Arkansas State	Sun Belt	Blake Anderson	825000.0	825000.0
10	Army	Ind.	Jeff Monken	932521.0	932521.0
22	Charlotte	C-USA	Brad Lambert	625000.0	625000.0
26	Colorado	Pac-12	Mike MacIntyre	2878500.0	2878500.0
54	Louisiana-Monroe	Sun Belt	Matt Viator	390000.0	390000.0
70	Navy	AAC	Ken Niumatalolo	2163000.0	2163000.0

	BonusPaid	Capacity	Rank	Revenue	Expenses	state	gsr	fgr
0	0.0	52237.0	57.0	60009782.0	54192115.0	CO	80.0	NaN
9	25000.0	30964.0	92.0	35550592.0	35550592.0	NaN	NaN	NaN
10	0.0	38000.0	NaN	NaN	NaN	NaN	NaN	NaN
22	0.0	NaN	76.0	40866430.0	38454232.0	NC	80.0	NaN
26	297000.0	53750.0	46.0	94935198.0	98413284.0	NaN	NaN	NaN
54	0.0	30427.0	178.0	15568952.0	16927856.0	NaN	NaN	NaN
70	0.0	34000.0	NaN	NaN	NaN	NaN	NaN	NaN

```
[105]: #Colleges missing Graduate Success Rate
coaches_comp[coaches_comp['gsr'].isna()]
```

```
[105]:
```

	School	Conference	Coach	SchoolPay	TotalPay	\
9	Arkansas State	Sun Belt	Blake Anderson	825000.0	825000.0	
10	Army	Ind.	Jeff Monken	932521.0	932521.0	
26	Colorado	Pac-12	Mike MacIntyre	2878500.0	2878500.0	
54	Louisiana-Monroe	Sun Belt	Matt Viator	390000.0	390000.0	
70	Navy	AAC	Ken Niumatalolo	2163000.0	2163000.0	

	BonusPaid	Capacity	Rank	Revenue	Expenses	state	gsr	fgr
9	25000.0	30964.0	92.0	35550592.0	35550592.0	NaN	NaN	NaN
10	0.0	38000.0	NaN	NaN	NaN	NaN	NaN	NaN
26	297000.0	53750.0	46.0	94935198.0	98413284.0	NaN	NaN	NaN
54	0.0	30427.0	178.0	15568952.0	16927856.0	NaN	NaN	NaN
70	0.0	34000.0	NaN	NaN	NaN	NaN	NaN	NaN

```
[124]: #Do a final clean of the coaches complete dataset
# Drop all schools with 0 TotalPay
coaches_comp = coaches_comp.drop(coaches_comp.index[coaches_comp.TotalPay == 0])
# Drop GSR null
coaches_comp = coaches_comp.dropna(axis=0, subset=['gsr'])
# Drop missing Stadium Capacity
coaches_comp = coaches_comp.dropna(axis=0, subset=['Capacity'])
# Drop missing Revenue
coaches_comp = coaches_comp.dropna(axis=0, subset=['Expenses'])
# Drop missing Expenses
coaches_comp = coaches_comp.dropna(axis=0, subset=['Revenue'])
print('The number of observations in the data set is:',len(coaches_comp))
```

The number of observations in the data set is: 98

##Data Exploration:

```
[125]: #Summarize the data
coaches_comp.describe()
```

```
[125]:
```

	SchoolPay	TotalPay	BonusPaid	Capacity	Revenue	\
count	9.800000e+01	9.800000e+01	9.800000e+01	98.000000	9.800000e+01	
mean	2.543377e+06	2.552000e+06	1.258214e+05	53311.846939	8.881076e+07	
std	1.993359e+06	1.998053e+06	2.311458e+05	24422.421562	5.132700e+07	
min	4.000000e+05	4.000000e+05	0.000000e+00	15000.000000	2.494847e+07	
25%	8.036280e+05	8.068875e+05	0.000000e+00	30371.250000	3.812054e+07	
50%	1.950004e+06	2.000000e+06	5.000000e+04	51521.000000	8.665144e+07	
75%	3.778438e+06	3.778438e+06	1.297500e+05	66139.000000	1.277516e+08	
max	8.307000e+06	8.307000e+06	1.350000e+06	107601.000000	2.238798e+08	

	Expenses	gsr	fgr
count	98	98	98
mean	1.993359e+06	0.000000e+00	0.000000e+00
std	1.993359e+06	0.000000e+00	0.000000e+00
min	0.000000e+00	0.000000e+00	0.000000e+00
25%	0.000000e+00	0.000000e+00	0.000000e+00
50%	0.000000e+00	0.000000e+00	0.000000e+00
75%	0.000000e+00	0.000000e+00	0.000000e+00
max	35550592.0	0.000000e+00	0.000000e+00

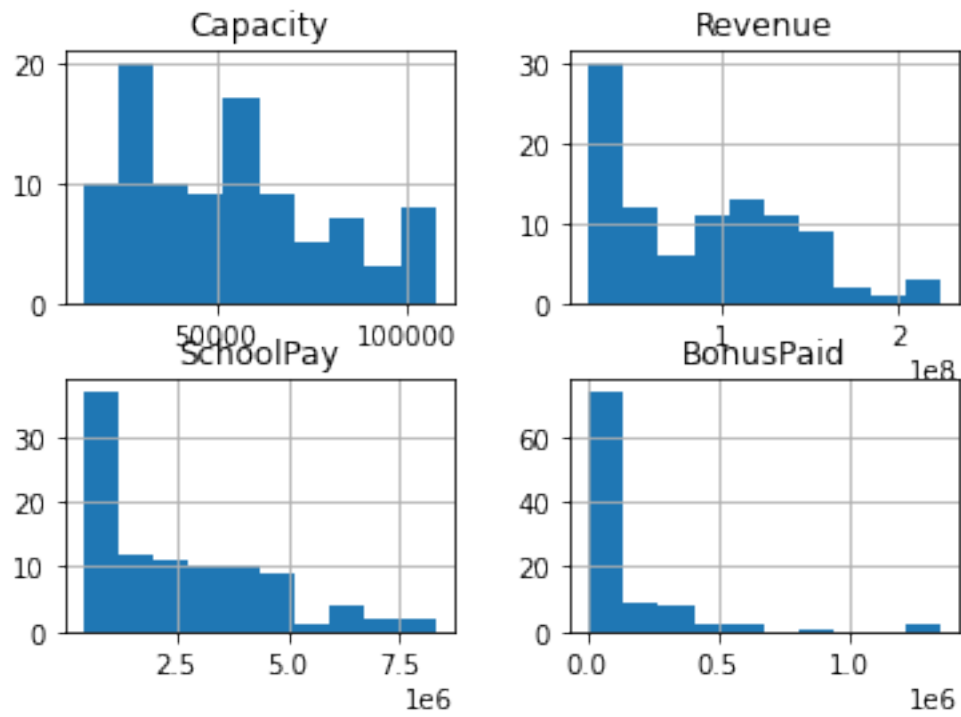
count	9.800000e+01	98.00000	97.000000
mean	8.680718e+07	78.22449	61.474227
std	4.879808e+07	8.34134	9.905904
min	2.491091e+07	54.00000	31.000000
25%	3.854889e+07	73.00000	55.000000
50%	8.298962e+07	78.00000	62.000000
75%	1.265413e+08	85.00000	69.000000
max	2.205730e+08	95.00000	84.000000

```
[126]: #Feature Exploration
df1 = coaches_comp[['SchoolPay', 'BonusPaid', 'Capacity', 'Revenue']]
df1.head(5)
```


```
[126]:   SchoolPay  BonusPaid  Capacity  Revenue
0   885000.0         0.0   52237.0  60009782.0
1   411000.0    50000.0   30000.0  37194485.0
2  8307000.0  500000.0  101821.0 164090889.0
4   712500.0  145000.0   24150.0  37996512.0
5  1600000.0         0.0   56037.0 105091389.0
```

```
[127]: #Plotting the features
vis1 = df1[['Capacity', 'Revenue', 'SchoolPay', 'BonusPaid']]
vis1.hist()
```


```
[127]: array([[<AxesSubplot:title={'center': 'Capacity'}>,
        <AxesSubplot:title={'center': 'Revenue'}>],
        [<AxesSubplot:title={'center': 'SchoolPay'}>,
        <AxesSubplot:title={'center': 'BonusPaid'}>]], dtype=object)
```




```
[128]: #Plot "Capacity" against "SchoolPay" to see relationship
plt.scatter(df1.Capacity, df1.SchoolPay, color='blue')
plt.xlabel("Capacity")
plt.ylabel("SchoolPay")
plt.show()
```



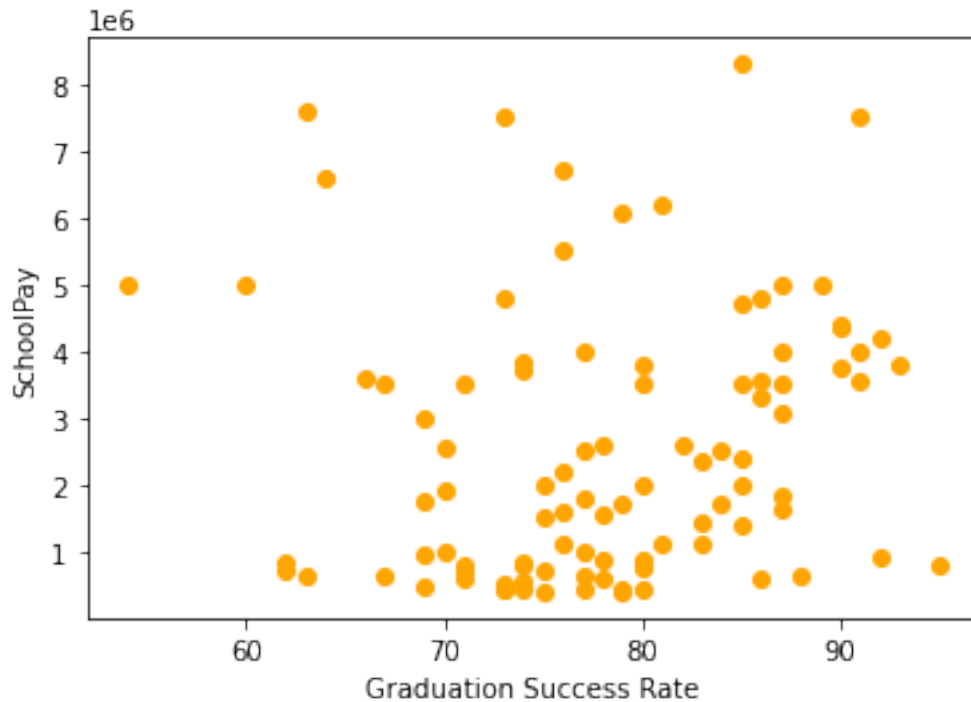
```
[129]: #Plot "Revenue" against "SchoolPay" to see relationship  
plt.scatter(df1.Revenue, df1.SchoolPay, color='red')  
plt.xlabel("Revenue")  
plt.ylabel("SchoolPay")  
plt.show()
```



```
[130]: #Plot "BonusPaid" against "SchoolPay" to see relationship  
plt.scatter(df1.BonusPaid, df1.SchoolPay, color='green')  
plt.xlabel("BonusPaid")  
plt.ylabel("SchoolPay")  
plt.show()
```



```
[174]: #Plot "gsr" against "SchoolPay" to see relationship
plt.scatter(coaches_comp.gsr, coaches_comp.SchoolPay, color='orange')
plt.xlabel("Graduation Success Rate")
plt.ylabel("SchoolPay")
plt.show()
```

##Regression Models:

```
[132]: #Create new df for first model
new_df = coaches_comp[['SchoolPay', 'BonusPaid', 'Capacity', 'Revenue', 'Expenses', 'gsr']]
```

```
[158]: #Create new df2 for second model
new_df2 = coaches_comp[['SchoolPay', 'Capacity', 'Revenue', 'Expenses', 'gsr']]
```

```
[133]: #Separate predictor variables in X
X=new_df[['BonusPaid', 'Capacity', 'Revenue', 'Expenses', 'gsr']]
```

```
[134]: #Put variable for prediction in Y
Y=new_df['SchoolPay']
```

```
[135]: X_train, X_test, y_train, y_test=train_test_split(X,Y,test_size=0.2)
```

```
[136]: regr = LinearRegression()
regr.fit(X_train,y_train)
```

```
[136]: LinearRegression()
```

```
[137]: #Print the coefficients and intercept
print ('Coefficients: ', regr.coef_)
```

```
print ('Intercept: ',regr.intercept_)
```

```
Coefficients: [ 1.04315387e+00  2.32735435e+01  2.02331230e-02  3.76234311e-03  
-2.39248569e+03]
```

```
Intercept: -761878.3721777261
```

```
##Prediction
```

```
[138]: y_pred=regr.predict(X_test)  
mse = mean_squared_error(y_test, y_pred)  
mae = mean_absolute_error(y_test, y_pred)  
score=r2_score(y_test, y_pred)  
#Display summary statistics for regression  
print("Mean absolute error : " + str(mae))  
print("Mean squared error : " + str(mse))  
print("r2_score : " + str(score))
```

```
Mean absolute error : 565661.8126893027
```

```
Mean squared error : 589858918954.0946
```

```
r2_score : 0.8542541751117876
```

```
##Model 2 - Fit model for SchoolPay with Capacity, Revenue, Expenses, and GSR
```

```
[159]: #A model predicting School pay based on stadium capacity, revenue, expenses, and gsr  
sp_model = str('SchoolPay ~ Capacity + Revenue + Expenses + gsr')  
  
# fit the model to the training set  
train_spmodel_fit = smf.ols(sp_model, data = new_df2).fit()  
  
# summary of model fit to the training set  
print(train_spmodel_fit.summary())
```

OLS Regression Results

```
=====
```

Dep. Variable:	SchoolPay	R-squared:	0.818
Model:	OLS	Adj. R-squared:	0.810
Method:	Least Squares	F-statistic:	104.6
Date:	Sun, 24 Jul 2022	Prob (F-statistic):	1.49e-33
Time:	19:28:00	Log-Likelihood:	-1476.6
No. Observations:	98	AIC:	2963.
Df Residuals:	93	BIC:	2976.
Df Model:	4		
Covariance Type:	nonrobust		

```
=====
```

	coef	std err	t	P> t	[0.025	0.975]
-----	-----	-----	-----	-----	-----	-----
Intercept	-1.158e+06	8.53e+05	-1.357	0.178	-2.85e+06	5.36e+05
Capacity	23.2531	8.543	2.722	0.008	6.287	40.219
Revenue	0.0183	0.011	1.630	0.107	-0.004	0.041

Expenses	0.0068	0.012	0.574	0.567	-0.017	0.030
gsr	3122.0315	1.07e+04	0.291	0.772	-1.82e+04	2.44e+04

Omnibus:	2.937	Durbin-Watson:	1.867
Prob(Omnibus):	0.230	Jarque-Bera (JB):	2.281
Skew:	0.301	Prob(JB):	0.320
Kurtosis:	3.442	Cond. No.	1.39e+09

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.39e+09. This might indicate that there are strong multicollinearity or other numerical problems.

##Recommended salary for Syracuse football head coach

```
[147]: #Syracuse Coach Salary
syr_sal = coaches_comp[coaches_comp['School'] == 'Syracuse']
regr.predict(syr_sal[['BonusPaid', 'Capacity', 'Revenue', 'Expenses', 'gsr']])

pred_syr_sal = regr.predict(syr_sal[['BonusPaid', 'Capacity', 'Revenue', '
    ↳'Expenses', 'gsr']])

# print amount
print("Syracuse Coach Pay should be $",(pred_syr_sal))
```

Syracuse Coach Pay should be \$ [2512146.28444074]

```
[170]: big10_sal = coaches_comp[coaches_comp['Conference'] == 'Big Ten']
print(big10_sal)
```

	School	Conference	Coach	SchoolPay	TotalPay \
43	Illinois	Big Ten	Lovie Smith	5000000.0	5000000.0
44	Indiana	Big Ten	Tom Allen	1830000.0	1830000.0
45	Iowa	Big Ten	Kirk Ferentz	4700000.0	4700000.0
58	Maryland	Big Ten	DJ Durkin	2512000.0	2512000.0
63	Michigan	Big Ten	Jim Harbaugh	7504000.0	7504000.0
64	Michigan State	Big Ten	Mark Dantonio	4390417.0	4390417.0
66	Minnesota	Big Ten	P.J. Fleck	3550000.0	3550000.0
71	Nebraska	Big Ten	Scott Frost	5000000.0	5000000.0
83	Ohio State	Big Ten	Urban Meyer	7600000.0	7600000.0
89	Penn State	Big Ten	James Franklin	4800000.0	4800000.0
91	Purdue	Big Ten	Jeff Brohm	3800000.0	3800000.0
93	Rutgers	Big Ten	Chris Ash	2200000.0	2200000.0
128	Wisconsin	Big Ten	Paul Chryst	3750000.0	3750000.0

	BonusPaid	Capacity	Rank	Revenue	Expenses	state	gsr	fg
--	-----------	----------	------	---------	----------	-------	-----	----

43	50000.0	60670.0	29.0	118565501.0	120168951.0	IL	87.0	73.0
44	0.0	52959.0	25.0	127832628.0	114822135.0	IN	87.0	70.0
45	600000.0	70585.0	14.0	151976026.0	146282275.0	IA	85.0	70.0
58	0.0	54000.0	33.0	108796303.0	108785924.0	MD	84.0	67.0
63	150000.0	107601.0	4.0	197820410.0	190952175.0	MI	91.0	72.0
64	125000.0	75005.0	18.0	140010865.0	135655740.0	MI	90.0	78.0
66	0.0	50805.0	24.0	130456454.0	129450256.0	MN	91.0	76.0
71	0.0	92000.0	21.0	136233460.0	124148206.0	NE	89.0	69.0
83	350000.0	104944.0	3.0	210548239.0	220572956.0	OH	63.0	45.0
89	300000.0	106572.0	6.0	164529326.0	160369805.0	PA	86.0	70.0
91	180000.0	62500.0	31.0	110844907.0	102026477.0	IN	80.0	64.0
93	50000.0	52454.0	39.0	103251280.0	103167344.0	NJ	76.0	55.0
128	290000.0	80321.0	11.0	157660107.0	154621828.0	WI	90.0	77.0

```
[172]: big10_sal = coaches_comp[coaches_comp['Conference'] == 'Big Ten']
regr.predict(big10_sal[['BonusPaid', 'Capacity', 'Revenue', 'Expenses', 'gsr']])

pred_big10_sal = regr.predict(big12_sal[['BonusPaid', 'Capacity', 'Revenue', 'Expenses', 'gsr']])

# print amount
print("Big 10 coach Pay should be $",(pred_big10_sal))
```

Big 10 coach Pay should be \$ [3345206.14525936 3280972.52211641 4928729.47937244 2904503.13512768 6402087.36025005 4242064.77713212 3329395.5712759 4389872.90866691 6984837.20114298 5757930.42580036 3315684.1929491 2806487.65777642 4966363.53920055]

```
[ ]:
```